

Notice (2/2)

- Python
 - Python 3.7 (<https://www.python.org/downloads/>)
 - opencv-contrib-python (3.4.2.17)
 - Matplotlib 3.1.1
 - UI framework: pyqt5 (5.15.1)

Assignment scoring (Total: 100%)

1. (20%) Image Processing (出題 : Mei)

1.1 (15%) Draw Contour

1.2 (5%) Count Rings

2. (20%) Camera Calibration (出題 : Jessica)

2.1 (4%) Corner detection

2.2 (4%) Find the intrinsic matrix

2.3 (4%) Find the extrinsic matrix

2.4 (4%) Find the distortion matrix

2.5 (4%) Show the undistorted result

3. (20%) Augmented Reality (出題 : Ming)

3.1 (10%) Show words on board

3.2 (10%) Show words vertically

4. (20%) Stereo Disparity Map (出題 : Maton)

4.1 (10%) Stereo Disparity Map

4.2 (10%) Checking the Disparity Value

5. (20%) Train a Cat-Dog Classifier Using ResNet50 (出題 : Benjamin)

5.1 (3%) Load the dataset and resize images

5.2 (3%) Plot class distribution of training dataset

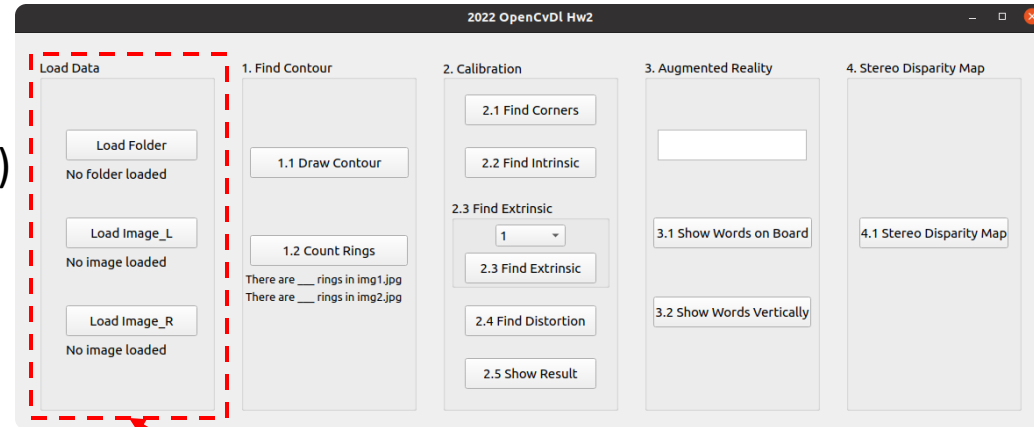
5.3 (3%) Show the structure of ResNet50 model

5.4 (3%) Set up 2 kinds of loss functions to train 2 ResNet50 models

5.5 (3%) Compare the accuracies of 2 ResNet50 models on validation dataset

5.6 (4%) Use the better-trained model to run inference and show the predicted class label

Question 5 needs to upload separately.



Don't fix your data path
(There is another dataset for demonstration)

5. Train a Cat-Dog Classifier Using ResNet50 (20%) (出題：Benjamin)

- 5.1 (3%) Load the dataset and resize images
- 5.2 (3%) Plot class distribution of training dataset
- 5.3 (3%) Show the structure of ResNet50 model
- 5.4 (3%) Set up 2 kinds of loss functions to train 2 ResNet50 models
- 5.5 (4%) Compare the accuracies of 2 ResNet50 models on validation dataset
- 5.6 (4%) Use the better-trained model to run inference and show the predicted class label

5.0 Train a Cat-Dog Classifier Using ResNet50

(出題 : Benjamin)

1. Objective

- 1) Learn how to train a ResNet50 model to **classify images** of cats and dogs
- 2) Learn how to deal with **imbalanced data**

2. ResNet50

- 1) Residual learning: avoid **degradation** problems
- 2) Bottleneck: build a **deeper** network without additional parameters

3. Cats and Dogs Dataset (modified from Kaggle Cats and Dogs Dataset)

- 1) Data type: JPG images
- 2) 2 classes: Cat and Dog
- 3) Datasets
 - (1) Path: /Download/Homework/Hw2/Dataset_OpenCvDI_Hw2_Q5.zip
 - (2) Training dataset: 16,200 JPG images in total. However, the dataset is **imbalanced**.
 - (3) Validation dataset: 1,800 JPG images in total.
 - (4) Inference dataset: 10 JPG images in total. It is for **testing the inference function** in your GUI program.

R. Reference

- 1) [Deep Residual Learning for Image Recognition](#)
- 2) [Kaggle Cats and Dogs Dataset](#)

Input: an RGB image
224x224x3c



112x112x64c

7x7 conv, 64/2

3x3 max pool/2 /2: stride = 2

56x56x256c

1x1 conv, 64
3x3 conv, 64
1x1 conv, 256

x3

Every conv layer is followed by a BN layer and a ReLU layer

28x28x512c

1x1 conv, 128/2
3x3 conv, 128
1x1 conv, 512

x3

BN: Batch Normalization
ReLU: Rectified Linear Unit

14x14x1024c

1x1 conv, 256/2
3x3 conv, 256
1x1 conv, 1024

x5

7x7x2048c

1x1 conv, 512/2
3x3 conv, 512
1x1 conv, 2048

x2

global avg pool

Output: FC, 1c FC: fully connected a confidence value (range: 0 ~ 1)

Figure: ResNet50 model structure




Figure: Residual Block (with bottleneck)

5.0 Train a Cat-Dog Classifier Using ResNet50

(出題： Benjamin)

4. Requirements

- 1) Train ResNet50 models with TensorFlow or PyTorch
- 2) Every chart should have a **chart title** and **axis titles**
- 3) Organize the files in this structure:

Hw2_05_StudentID_Name_Version // project folder

| -- model // folder to put trained models

| -- inference_dataset

| -- Cat

| -- Dog

| -- main.py // codes for your GUI program

| -- train.py // codes for model training

| -- : // other files or folders you need

Notice: Please include the inference dataset in your homework file.

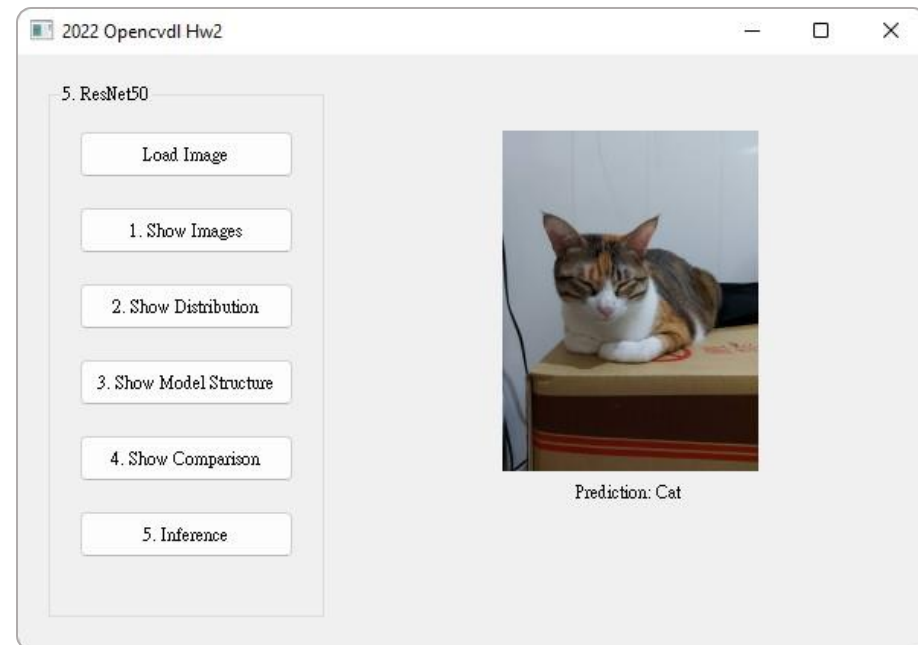


Figure: GUI example

5.1 (3%) Load the dataset and resize images

(出題 : Benjamin)

1. Objective

1) At home:

- (1) Load the training dataset and validation dataset

→ Hint:

(a) TensorFlow ([tutorial](#)): `tf.keras.utils.image_dataset_from_directory()`

(b) PyTorch ([tutorial](#)): `torch.utils.data.Dataset`

- (2) Resize images to **224×224×3c** (RGB)

→ Hint:

(a) TensorFlow: `tf.keras.utils.image_dataset_from_directory(image_size=(?, ?))`

(b) PyTorch ([tutorial](#)): `torchvision.transform`

2) When the demo:

- (1) Click the button “1. Show Images”

- (2) Load the **inference dataset**

- (3) Resize images to **224×224×3c** (RGB)

- (4) Get 1 image from **each class** in the inference dataset

- (5) Show images in a **new window**

→ Hint: use `matplotlib.pyplot` functions to show images ([tutorial](#)):

(a) `figure()`

(b) `imshow()`

(c) `subplot()`

(d) `title()`

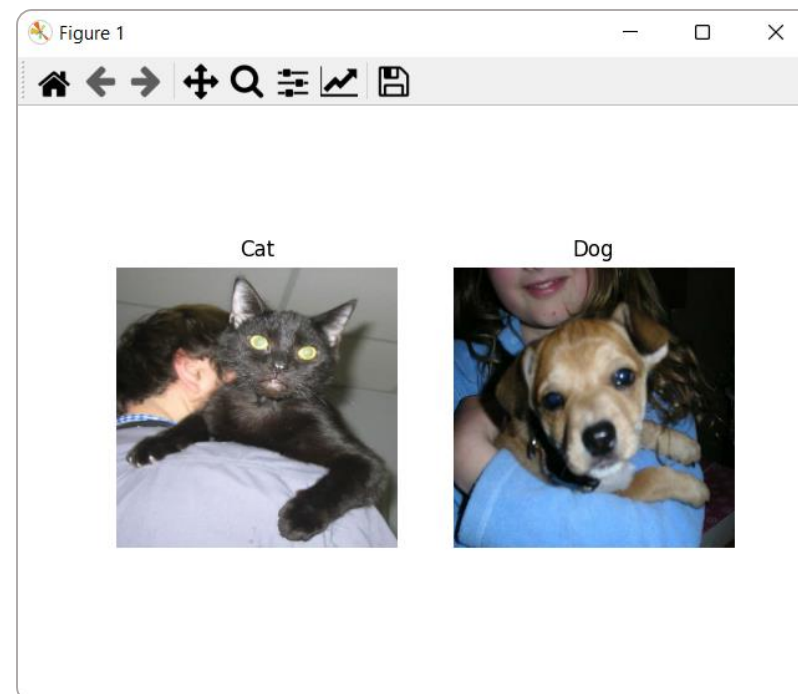
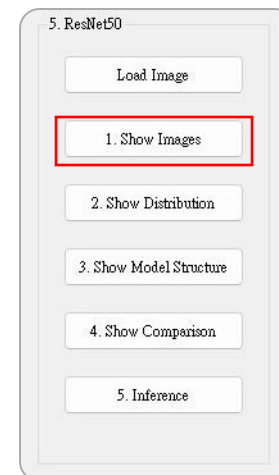


Figure: 1 image from each class

Notice: this is an example, the images might differ

5.2 (3%) Plot class distribution of training dataset (出題：Benjamin)

1. Objective

1) At home:

- (1) Load the training dataset
- (2) Iterate through the dataset and count **the number of images of each class**
- (3) Plot the class distribution

→ Hint: use matplotlib.pyplot functions to plot

- (a) figure()
- (b) bar()
- (c) xticks()
- (d) ylabel()

(4) **Save the figure**

2) When the demo:

- (1) Click the button **"2. Show Distribution"**
- (2) Show the **saved figure** of class distribution in a new window

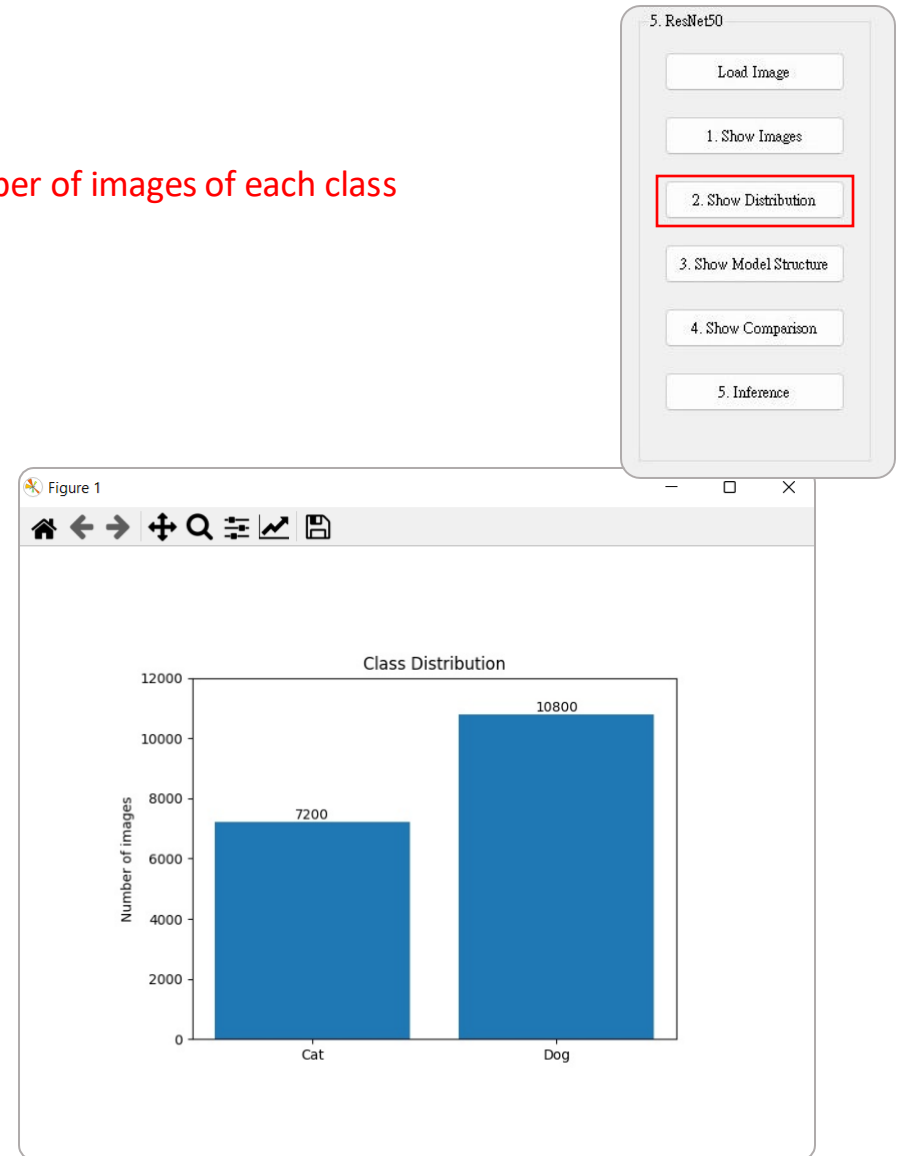


Figure: Class Distribution

Notice: this is an example, the numbers might differ

5.3 (3%) Show the structure of ResNet50 model

(出題 : Benjamin)

1. Objective

1) At home:

(1) Build a ResNet50 model

→ Hint:

(a) TensorFlow: `tf.keras.applications.resnet50.ResNet50()`

(b) PyTorch: `torchvision.models.resnet50()`

(2) Replace the output layer to a FC (Fully Connected) layer of **1 node** with a **Sigmoid** activation function

→ Hint:

(a) TensorFlow ([tutorial](#)): `tf.keras.layers.Dense(1, activation='sigmoid')`

(b) PyTorch ([tutorial](#)): `torch.nn.Linear(2048, 1), torch.nn.Sigmoid`

If the class label of Cat is 1, the output value (range: 0 ~ 1) should be **close to 1** for cat images, and vice versa.

(3) Run the function to show the structure **in the terminal**

→ Hint:

(a) TensorFlow: `tf.keras.Model.summary()`

(b) PyTorch: `torchsummary`

2) When the demo:

(1) Click the button **"3. Show Model Structure"**

(2) Run the function to show the structure **in the terminal**



Layer (type)	Feature map shape (B, H, W, C)	Num. of param.	Layer connecting to
input_1 (InputLayer)	(None, 224, 224, 3)	0	[]
conv1_pad (ZeroPadding2D)	(None, 230, 230, 3)	0	['input_1[0][0]']
conv1_conv (Conv2D)	(None, 112, 112, 64)	9472	['conv1_pad[0][0]']
conv1_bn (BatchNormalization)	(None, 112, 112, 64)	256	['conv1_conv[0][0]']
conv1_relu (Activation)	(None, 112, 112, 64)	0	['conv1_bn[0][0]']
pool1_pad (ZeroPadding2D)	(None, 114, 114, 64)	0	['conv1_relu[0][0]']
pool1_pool (MaxPooling2D)	(None, 56, 56, 64)	0	['pool1_pad[0][0]']
conv2_block1_1_conv (Conv2D)	(None, 56, 56, 64)	4160	['pool1_pool[0][0]']
conv2_block1_1_bn (BatchNormalization)	(None, 56, 56, 64)	256	['conv2_block1_1_conv[0][0]']
conv2_block1_1_relu (Activation)	(None, 56, 56, 64)	0	['conv2_block1_1_bn[0][0]']

Figure: the structure of ResNet50 model

5.3 (3%) Show the structure of ResNet50 model

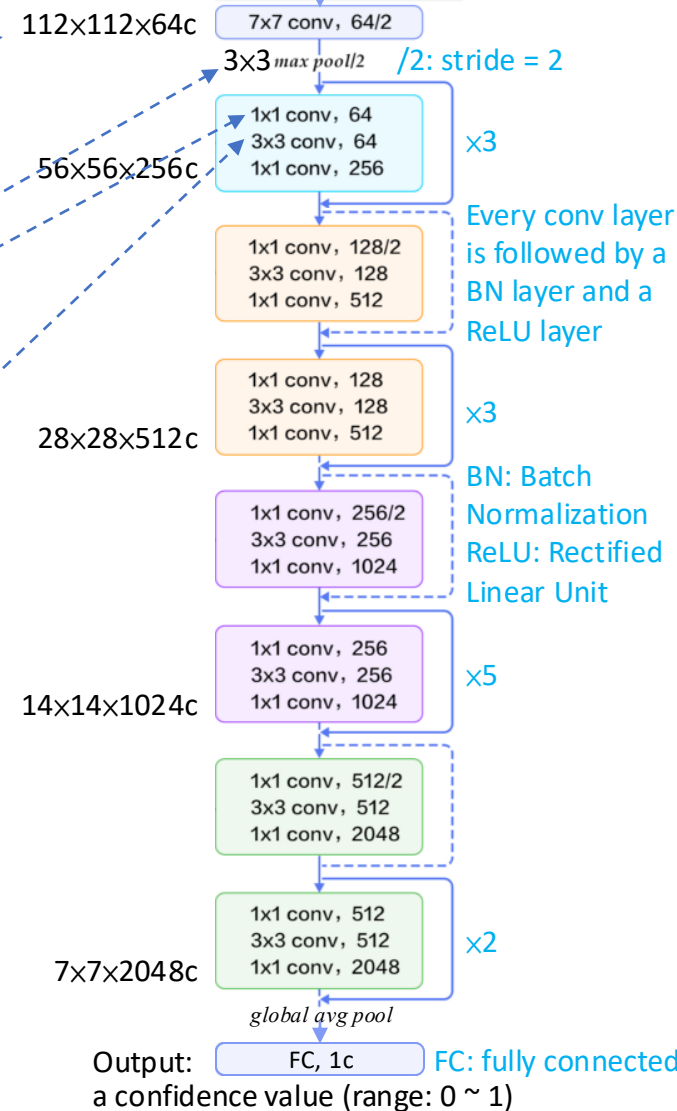
(出題：Benjamin)

2. How to read the structure

Layer (type)	Feature map shape (B, H, W, C)	Num. of param.	Layer connecting to
input_1 (InputLayer)	(None, 224, 224, 3)	0	conv1_pad
conv1_pad (ZeroPadding2D)	(None, 230, 230, 3)	0	conv1_conv
conv1_conv (Conv2D)	(None, 112, 112, 64)	9472	conv1_bn
conv1_bn (BatchNormalization)	(None, 112, 112, 64)	256	conv1_relu
conv1_relu (Activation)	(None, 112, 112, 64)	0	pool1_pad
pool1_pad (ZeroPadding2D)	(None, 114, 114, 64)	0	pool1_pool
pool1_pool (MaxPooling2D)	(None, 56, 56, 64)	0	conv2_block1_1_conv
conv2_block1_1_conv (Conv2D)	(None, 56, 56, 64)	4160	conv2_block1_1_bn
conv2_block1_1_bn (BatchNormalization)	(None, 56, 56, 64)	256	conv2_block1_1_relu
conv2_block1_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block1_2_conv
conv2_block1_2_conv (Conv2D)	(None, 56, 56, 64)	36928	conv2_block1_2_bn
conv2_block1_2_bn (BatchNormalization)	(None, 56, 56, 64)	256	conv2_block1_2_relu
conv2_block1_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block1_0_conv
conv2_block1_0_conv (Conv2D)	(None, 56, 56, 256)	16640	conv2_block1_3_conv
conv2_block1_3_conv (Conv2D)	(None, 56, 56, 256)	16640	conv2_block1_0_bn
conv2_block1_0_bn (BatchNormalization)	(None, 56, 56, 256)	1024	conv2_block1_3_bn
conv2_block1_3_bn (BatchNormalization)	(None, 56, 56, 256)	1024	

Figure: the structure of ResNet50 model

Input: an RGB image
224×224×3c



5.4 (3%) Set up 2 kinds of loss functions to train 2 ResNet50 models

(出題 : Benjamin)

1. Objective

1) At home:

- (1) Set up **Focal Loss (α -balanced)** and **Binary Cross Entropy** in codes for model training (train.py)

→ Hint:

(a) TensorFlow: `tfa.losses.SigmoidFocalCrossEntropy()`, `tf.keras.losses.BinaryCrossEntropy()`

(b) PyTorch: `torchvision.ops.sigmoid_focal_loss()`, `torch.nn.function.binary_cross_entropy()`

Because you have to use **Focal Loss**, which accepts 1 confidence value, you may replace the output layer of ResNet50 model to a FC layer of **1 node** with a **Sigmoid** activation function.

- (2) Train **2 ResNet50 models** with training dataset using **2 loss functions**

→ Hint:

(a) TensorFlow: `tf.keras.Model.fit()`

(b) PyTorch ([tutorial](#)): write a for loop to train the model

2) When the demo:

- (1) Show your **codes** about loss functions in train.py

```
60  ### Q5.4 Set up 2 kinds of loss functions to train 2 ResNet50 models
61  # 1st loss function
62  loss_function = tfa.losses.SigmoidFocalCrossEntropy(alpha=0.4, gamma=1.0)
63
64  # 2nd loss function
65  # loss_function = tf.keras.losses.BinaryCrossentropy()
66
67  optimizer = tf.keras.optimizers.Adam(learning_rate=8e-5)
68
69  # configure loss function and optimizer for training
70  model.compile(optimizer=optimizer, loss=loss_function, metrics=['accuracy'])
```

Train a model with Focal Loss

Train another model with Binary Cross Entropy

Figure: Code example (you can follow the hyperparameter settings)

5.4 (3%) Set up 2 kinds of loss functions to train 2 ResNet50 models

(出題 : Benjamin)

2. What is Focal Loss?

- $FL(p_t) = -\alpha_t(1 - p_t)^\gamma \log(p_t)$
 - 1) $p_t = \begin{cases} p & \text{if } y = 1 \\ 1 - p & \text{otherwise,} \end{cases}$
 - 2) y : **ground-truth** class label (0 or 1)
 - 3) p : model's **output confidence** value for the input with class label $y = 1$ (range: 0 ~ 1)
 - 4) $\alpha_t = \begin{cases} \alpha & \text{if } y = 1 \\ 1 - \alpha & \text{otherwise,} \end{cases}$
 - 5) α : class imbalance weighting factor (Usually set to $1 - \text{frequency of } y = 1 \text{ examples}$)
 - 6) γ : focusing parameter (Usually set to 2)
- Focal Loss mitigates **class imbalance** by assigning **more weights** to hard examples and less weights to easy examples.
- Focal Loss is for **binary classification**, ex: Foreground and Background, Face and Non-Face, or Cat and Dog.
- Reference: [Focal Loss for Dense Object Detection](#)

Ex: we have

- 1) 30 examples with $y = 0$
 - 2) 70 examples with $y = 1$
- Frequency of $y = 1$ examples is $\frac{70}{100} = 0.7$
- $\alpha = 1 - 0.7 = 0.3$

5.5 (4%) Compare the accuracies of 2 ResNet50 models on validation dataset

(出題：Benjamin)

1. Objective

1) At home:

(1) Validate 2 ResNet50 models with validation dataset

→ Hint:

(a) TensorFlow: `tf.keras.Model.evaluate()`

(b) PyTorch ([tutorial](#)): write a for loop to validate the model

(2) Plot the accuracy values with a bar chart

(3) Save the figure

2) When the demo:

(1) Click the button “4. Show Comparison”

(2) Show the saved figure of class distribution in a new window

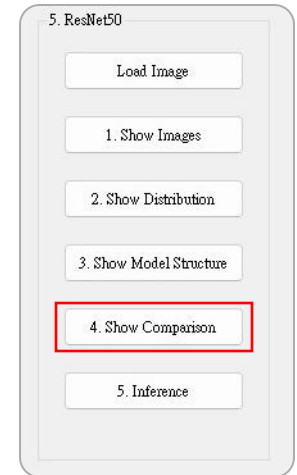


Figure: Accuracy Comparison

Notice: this is an example, the numbers might differ

5.6 (4%) Use the better-trained model to run inference and show the predicted class label

(出題 : Benjamin)

1. Objective

1) At home:

(1) Load the trained model

→ Hint:

(a) TensorFlow: `tf.keras.Model.load_weights()`

(b) PyTorch: `torch.nn.Module.load_state_dict()`

(2) Click the button “Load Image” to select 1 image arbitrarily

→ Hint: `PyQt5.QtWidgets.QFileDialog.getOpenFileName()`

(3) Show the loaded image in the GUI

(4) Resize the loaded image to 224×224×3c (RGB)

(5) Click the button “5. Inference” to run inference on the resized image

→ Hint:

(a) TensorFlow: `tf.keras.Model.predict_step()`

(b) PyTorch: pass an image when calling `torch.nn.M` object to run inference, ex: `trained_model(img)`

(6) Show the predicted class label

→ Hint: decide the class label with a threshold of the output value.

Ex: class label = $\begin{cases} \text{Cat, } output < thresh \\ \text{Dog, } output \geq thresh \end{cases}$
 $thresh = 0.5$

2) When the demo: repeat the process

(2) Select 1 image arbitrarily

